

About the presentation

- We are in 370 during each session. Be there on time! We have pizzas. :)
- Check piazza for the rubric and sign up. There are still some empty slots.
- **Imagine what it is like for your audience!**
- Don't write too much text on the slides - noone reads it. Write short bullet points.
- Figures must have x and y labels, the characters must be readable from a distance, the objects of the graph must be visible.
- EDA figures:
 - three different visualization types!
 - e.g., one bar plot, one scatter plot, one heatmap.
 - if you only show bar plots, those are all of the same type and points will be subtracted.

Mud card

- **Pls give more exercises on coordinates and p-criti**
 - in HW5 :)
- **What does accuracy refer to? the accuracy of the model?**
 - accuracy as the classification metric $(TP + TN) / (TP + TN + FP + FN)$
 - accuracy as the predictive power of the model ("the model is accurate" meaning that it gives good predictions)
- **Is there a reason why you would change the value for epsilon in log loss?**
 - $1e-15$ is used because python's float contains 16 significant digits (double precision)
 - $1e-15$ has at least one significant digit
 - if you represent your numbers with a different precision level, you might want to adjust epsilon accordingly
 - e.g., single precision (8 significant digits), $\text{eps} = 1e-7$
- **It would be great if you could clarify logloss once again in the next class**
- **What is y_{true} in log loss**
- **Can you explain extreme cases in logloss?**

The logloss metric

$$\text{logloss} = -\frac{1}{N} \sum (y_{\text{true}} \ln(p_{\text{pred}}) + (1 - y_{\text{true}})(1 - \ln(1 - p_{\text{pred}})))$$

- p_{pred} is the predicted probability of the **positive class**
- the predicted probabilities are not converted into predicted classes
- two scenarios:
 - $y_{\text{true}} = 0$ - left term disappears
 - $y_{\text{true}} = 1$ - right term disappears
- $\log(0)$ is undefined
 - p_{pred} is replaced with $\max(\min(p, 1 - 10^{-15}), 10^{-15})$ to avoid this issue

The extreme cases

- the classifier is confidently wrong
 - $p_{\text{pred}} = 10^{-15}$ for points in class 1
 - $p_{\text{pred}} = 1 - 10^{-15}$ for points in class 0

$$\text{logloss} = -\frac{1}{N} \sum \ln(10^{-15}) = -\ln(10^{-15})$$

$$\text{logloss} \sim 34.5$$

- the classifier is correct
 - $p_{\text{pred}} = 10^{-15}$ for points in class 0
 - $p_{\text{pred}} = 1 - 10^{-15}$ for points in class 1

$$\text{logloss} = -\frac{1}{N} \sum (1 - 0)(1 - \ln(1 - 10^{-15})) = 10^{-15} \text{ for class 0}$$

$$\text{logloss} = -\frac{1}{N} \sum 1 * \ln(1 - 10^{-15}) = 10^{-15} \text{ for class 1}$$

$$\text{logloss} \sim 0$$

```
In [1]: from sklearn.metrics import log_loss
import pandas as pd
import numpy as np
df = pd.read_csv('data/true_labels_pred_probs.csv')

y_true = df['y_true']
pred_prob_class1 = df['pred_prob_class1']

print(log_loss(y_true, pred_prob_class1))
help(log_loss)
```

0.35015190545328556

Help on function log_loss in module sklearn.metrics.classification:

```
log_loss(y_true, y_pred, eps=1e-15, normalize=True, sample_weight=None,
labels=None)
```

Log loss, aka logistic loss or cross-entropy loss.

This is the loss function used in (multinomial) logistic regression and extensions of it such as neural networks, defined as the negative log-likelihood of the true labels given a probabilistic classifier's

predictions. The log loss is only defined for two or more labels. For a single sample with true label y_t in $\{0,1\}$ and estimated probability y_p that $y_t = 1$, the log loss is

$$-\log P(y_t|y_p) = -(y_t \log(y_p) + (1 - y_t) \log(1 - y_p))$$

Read more in the :ref:`User Guide <log_loss>`.

Parameters

y_true : array-like or label indicator matrix
Ground truth (correct) labels for n_{samples} samples.

y_pred : array-like of float, shape = (n_{samples} , n_{classes}) or (n_{samples} ,

Predicted probabilities, as returned by a classifier's `predict_proba` method. If `y_pred.shape = (n_samples,)` the probabilities provided are assumed to be that of the positive class. The labels in `y_pred` are assumed to be ordered alphabetically, as done by `:class:`preprocessing.LabelBinarizer``.

eps : float

Log loss is undefined for $p=0$ or $p=1$, so probabilities are clipped to $\max(\text{eps}, \min(1 - \text{eps}, p))$.

normalize : bool, optional (default=True)

If true, return the mean loss per sample. Otherwise, return the sum of the per-sample losses.

sample_weight : array-like of shape = [n_{samples}], optional
Sample weights.

labels : array-like, optional (default=None)

If not provided, labels will be inferred from `y_true`. If `labels`

is `None` and `y_pred` has shape (n_{samples} ,) the labels are

assumed to be binary and are inferred from `y_true`.

.. versionadded:: 0.18

Returns

loss : float

Examples

```
>>> from sklearn.metrics import log_loss
>>> log_loss(["spam", "ham", "ham", "spam"], # doctest: +ELLIPSIS
...         [[.1, .9], [.9, .1], [.8, .2], [.35, .65]])
0.21616...
```

References

C.M. Bishop (2006). Pattern Recognition and Machine Learning. Springer,
p. 209.

Notes

The logarithm used is the natural logarithm (base-e).

Evaluation metrics in regression and gradient descent

By the end of this lecture, you will be able to

- Describe the most commonly used regression metrics
- Describe what the cost function is
- Explain how a simple gradient descent algorithm works

Evaluation metrics in regression and gradient descent

By the end of this lecture, you will be able to

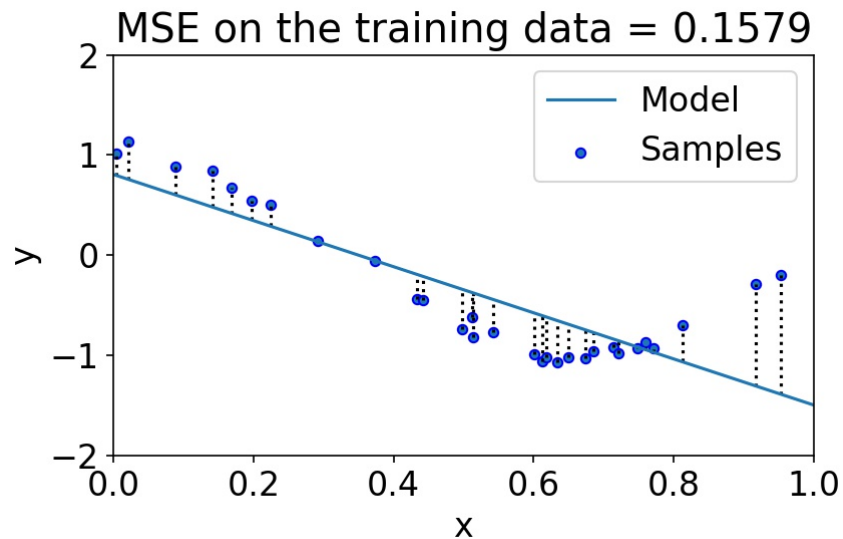
- **Describe the most commonly used regression metrics**
- Describe what the cost function is
- Explain how a simple gradient descent algorithm works

Regression metrics

- the target variable is continuous
- the predicted values are also continuous
- regression metrics measure some type of difference between y (true values) and y' (predicted values)
- three types of metrics:
 - unit of metric is different than the unit of target variable
 - same units
 - dimensionless

Mean Squared Error

$$MSE(y, y') = \frac{1}{n} \sum_{i=1}^n (y_i - y'_i)^2$$



Root Mean Square Error

$$RMSE(y, y') = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - y'_i)^2}$$

Mean Absolute Error

$$MAE(y, y') = \frac{1}{n} \sum_{i=1}^n |y_i - y'_i|$$

Both RMSE and MAE have the same unit as the target variable.

R2 score - coefficient of determination

$$R^2(y, y') = 1 - \frac{\sum_{i=1}^n (y_i - y'_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2},$$

where \bar{y} is the mean of y .

- $R^2 = 1$ is the perfect regression model ($y == y'$)
- $R^2 = 0$ is as good as a constant model that always predicts the expected value of y , \bar{y}
- $R^2 < 0$ is a bad regression model

R2 is dimensionless.

```
In [2]: from sklearn.metrics import mean_squared_error
        from sklearn.metrics import mean_absolute_error
        from sklearn.metrics import r2_score
```

- RMSE is not implemented in sklearn, but you can calculate it as
`np.sqrt(mean_squared_error(y_true, y_pred))`
- you can find more on regression metrics [here \(https://scikit-learn.org/stable/modules/model_evaluation.html#regression-metrics\)](https://scikit-learn.org/stable/modules/model_evaluation.html#regression-metrics).

Exercise 1

Read in `data/reg_preds.csv` . It contains two columns:

- `y_true`: value of owner-occupied homes in \$1000's in Boston
- `y_pred`: predictions of a regression model

What's the ratio between the MSE and the variance of the home values? How does this ratio relate to the R^2 score?

```
In [ ]:
```

Evaluation metrics in regression and gradient descent

By the end of this lecture, you will be able to

- Describe the most commonly used regression metrics
- **Describe what the cost function is**
- Explain how a simple gradient descent algorithm works

Supervised ML algorithms

- What mathematical model is used to represent the data? What are the parameters of the model?
- How do we compare different parameter values of the same model?
 - this is done with the cost function
 - WARNING!
 - the evaluation metric is used to compare different models!
 - the cost function compares different parameter values of the same model!
- What algorithm do we use to find the best model parameter values?
 - e.g., brute force, gradient descent, backpropagation

Today

- The mathematical model is linear regression: ###

$$y'_i = \theta_0 + x_{i1}\theta_1 + x_{i2}\theta_2 + \dots = \theta_0 + \sum_{j=1}^m \theta_j x_{ij}$$
 where y'_i is the prediction of the linear regression model and θ are parameters.
- The cost function is MSE
- We will find the best parameter values by brute force first, then simple gradient descent.

Let's generate some data

```
In [6]: # load packages and generate data
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import matplotlib

matplotlib.rcParams.update({'font.size': 16})

# the true function to generate y (gaussian noise is added later)
def true_fun(X):
    return np.cos(1.5 * np.pi * X)

# fix the seed so code is reproducible
np.random.seed(10)

# generate n_samples points
n_samples = 30

# generate data
X = np.random.rand(n_samples)
y = true_fun(X) + np.random.randn(n_samples) * 0.1 # noise added here
```



```
In [7]: df = pd.DataFrame()
for i in range(10):
    df['x'+str(i+1)] = X**(i+1)
df['y'] = y
print(df.head())
df.to_csv('data/regression_example.csv', index=False)
```

	x1	x2	x3	x4	x5	
x6 \						
0	0.771321	0.594936	0.458886	3.539483e-01	2.730076e-01	2.105764e-01
1	0.020752	0.000431	0.000009	1.854537e-07	3.848527e-09	7.986443e-11
2	0.633648	0.401510	0.254416	1.612103e-01	1.021507e-01	6.472758e-02
3	0.748804	0.560707	0.419860	3.143926e-01	2.354184e-01	1.762822e-01
4	0.498507	0.248509	0.123884	6.175684e-02	3.078622e-02	1.534715e-02

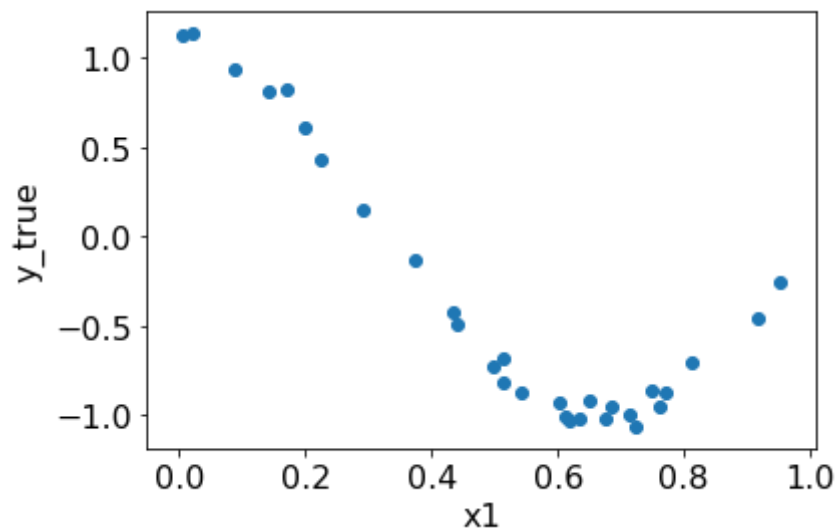
	x7	x8	x9	x10	y
0	1.624219e-01	1.252794e-01	9.663058e-02	7.453316e-02	-0.870922
1	1.657343e-12	3.439309e-14	7.137237e-16	1.481116e-17	1.135022
2	4.101452e-02	2.598878e-02	1.646774e-02	1.043476e-02	-1.015044
3	1.320008e-01	9.884272e-02	7.401381e-02	5.542183e-02	-0.864701
4	7.650660e-03	3.813908e-03	1.901260e-03	9.477913e-04	-0.728846

```
In [8]: def predict(X,theta):
    if len(np.shape(theta)) != 2:
        theta = np.array(theta)[np.newaxis,:] # just a numpy trick to make the dot product work
    y_pred = theta[0,0] + X.dot(theta[0,1:]) # intercept + theta_i*x_i
    return y_pred

def cost_function(X,y_true,theta):
    """
    Take in a numpy array X,y_true, theta and generate the cost function of using theta as parameter in a linear regression model
    """
    m = len(y)
    theta = np.array(theta)[np.newaxis,:] # just a numpy trick to make the dot product work in predict
    y_pred = predict(X,theta)
    cost = (1/m) * np.sum(np.square(y_true-y_pred)) # this is MSE
    return cost
```

For simplicity, let's focus on x1 and y only!

```
In [9]: plt.scatter(df['x1'],df['y'])  
plt.xlabel('x1')  
plt.ylabel('y_true')  
plt.savefig('figures/data.png',dpi=300)  
plt.show()
```



$$y'_i = \theta_0 + x_{i1}\theta_1$$

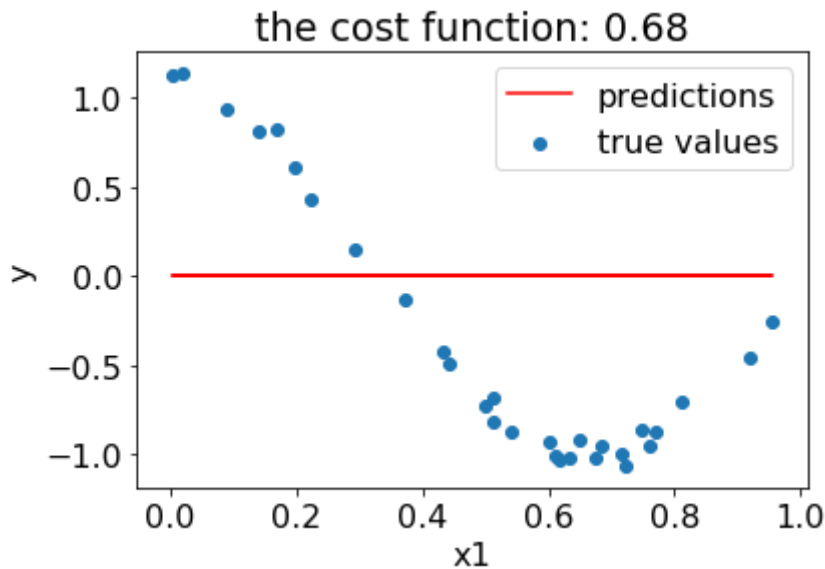
- θ_0 is the intercept
- θ_1 is the slope

We are looking for the best fit line!

For a given θ vector, the cost function returns the MSE.

```
In [11]: theta = [0,0] # intercept is theta[0], the slope is theta[1]

plt.scatter(df['x1'],df['y'],label='true values')
plt.plot(df['x1'],predict(df['x1'].values[:,np.newaxis],theta),label='pr
edictions',color='r')
plt.title('the cost function: '+str(np.around(cost_function(df['x1'].val
ues[:,np.newaxis],df['y'],theta),2)))
plt.xlabel('x1')
plt.ylabel('y')
plt.legend()
plt.savefig('figures/line_fit.png',dpi=300)
plt.show()
```



What we want:

- Find the theta vector that minimizes the cost function!
 - that's our best fit model

How we do it:

- brute force
 - create a grid of theta[0] and theta[1] values
 - loop through all theta vectors on the grid
 - find the theta vector that comes with the smallest cost

```
In [12]: n_vals = 101

theta0 = np.linspace(-5,5,n_vals) # the intercept values to explore
theta1 = np.linspace(-5,5,n_vals) # the slope values to explore

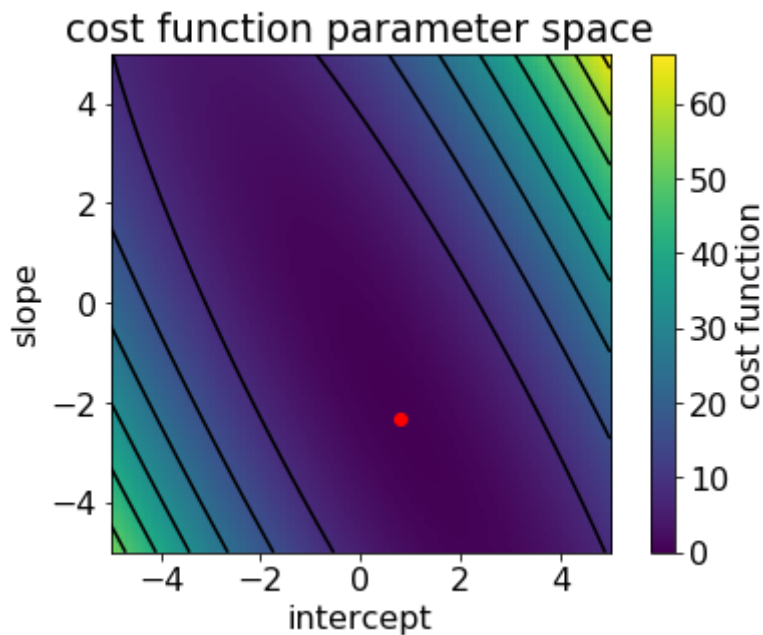
cost = np.zeros([len(theta0),len(theta1)]) # the cost function's value f
or each theta

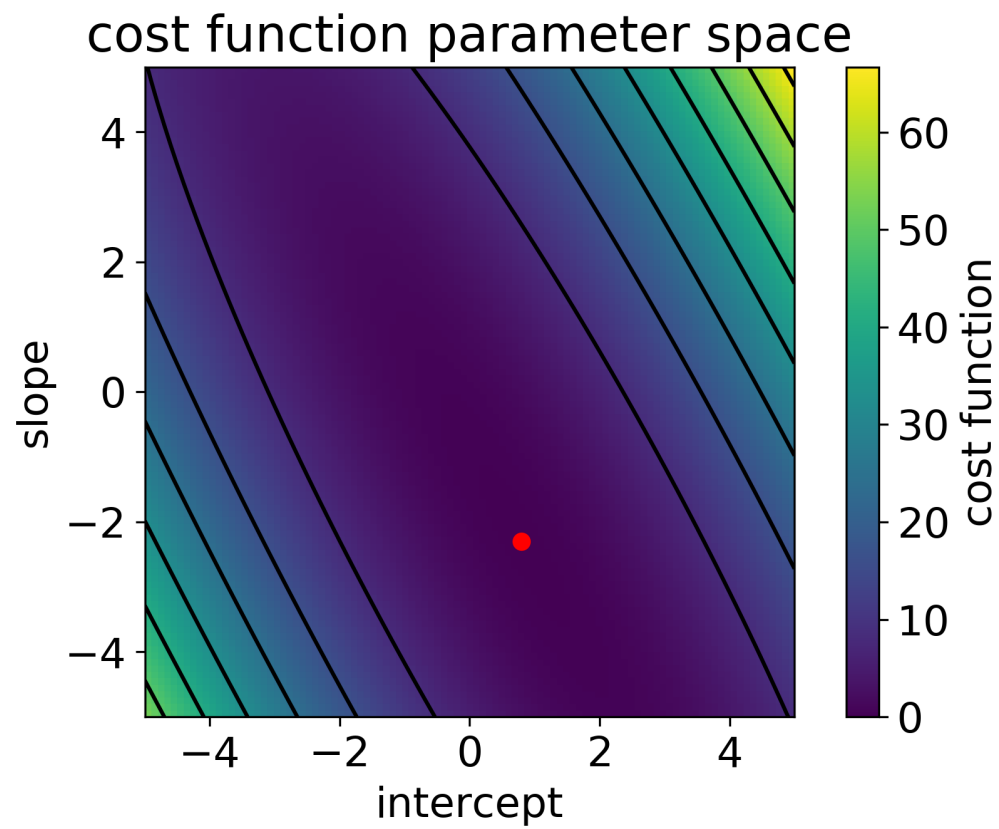
# loop through all intercept-slope combinations and calculate the cost f
unction
for i in range(n_vals):
    for j in range(n_vals):
        theta = [theta0[i],theta1[j]]
        cost[i,j] = cost_function(df['x1'].values[:,np.newaxis],df['y'],
theta)

print('min(cost):',np.min(cost))
min_coords = np.unravel_index(cost.argmin(),np.shape(cost))
print('best intercept:',theta0[min_coords[0]])
print('best slope:',theta1[min_coords[1]])

min(cost): 0.14851645747269254
best intercept: 0.8000000000000007
best slope: -2.3
```

```
In [13]: plt.figure(figsize=(6.4,4.8))
ax = plt.gca()
extent = (np.min(theta0),np.max(theta0),np.min(theta1),np.max(theta1))
fig = ax.imshow(cost.T,origin='lower',extent=extent,vmin=0)
plt.colorbar(fig,label='cost function')
ax.contour(theta0,theta1,cost.T,levels=10,colors='black')
plt.scatter(theta0[min_coords[0]],theta1[min_coords[1]],c='r')
ax.xaxis.set_ticks_position("bottom")
plt.xlabel('intercept')
plt.ylabel('slope')
plt.title('cost function parameter space')
plt.tight_layout()
plt.savefig('figures/cost_function.png',dpi=300)
plt.show()
```





The brute force approach works but...

- the number of theta vectors to loop through explodes with the number of features we have
 - with n features, we would need to loop through $\sim 100^n$ theta vectors.
 - no guarantee that the best theta vector is within our grid.
- We need to use a smarter numerical method to find the best theta!
 - gradient descent to the rescue!

Evaluation metrics in regression and gradient descent

By the end of this lecture, you will be able to

- Describe the most commonly used regression metrics
- Describe what the cost function is
- **Explain how a simple gradient descent algorithm works**

```
In [14]: def gradient_descent(X,y_true,theta,learning_rate=0.01,iterations=100):
    '''
    X      = Matrix of X with added bias units
    y      = Vector of Y
    theta=Vector of thetas np.random.randn(j,1)
    learning_rate
    iterations = no of iterations

    Returns the final theta vector and array of cost history over no of
    iterations
    '''
    m = len(y_true)
    theta = np.array(theta)[np.newaxis,:]

    cost_history = np.zeros(iterations)
    theta_history = np.zeros([iterations,np.shape(theta)[1]])
    for it in range(iterations):

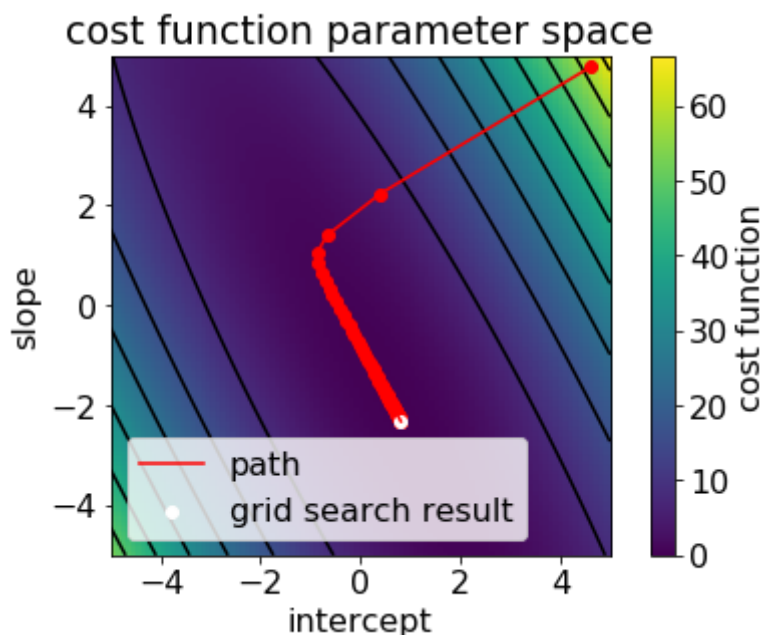
        y_pred = predict(X,theta)
        delta_theta = np.zeros(np.shape(theta)) # the step we take
        # the derivative of the cost function with respect to the interc
ept
        delta_theta[0,0] = (1/m) * sum(y_pred - y_true) * learning_rate
        # the derivative of the cost function with respect to the slopes
* learning_rate
        delta_theta[0,1:] = (1/m)*learning_rate*( X.T.dot((y_pred - y_tr
ue)))
        theta = theta - delta_theta # update theta so we move down the g
radient
        theta_history[it] = theta[0]
        cost_history[it] = cost_function(X,y_true,theta[0])

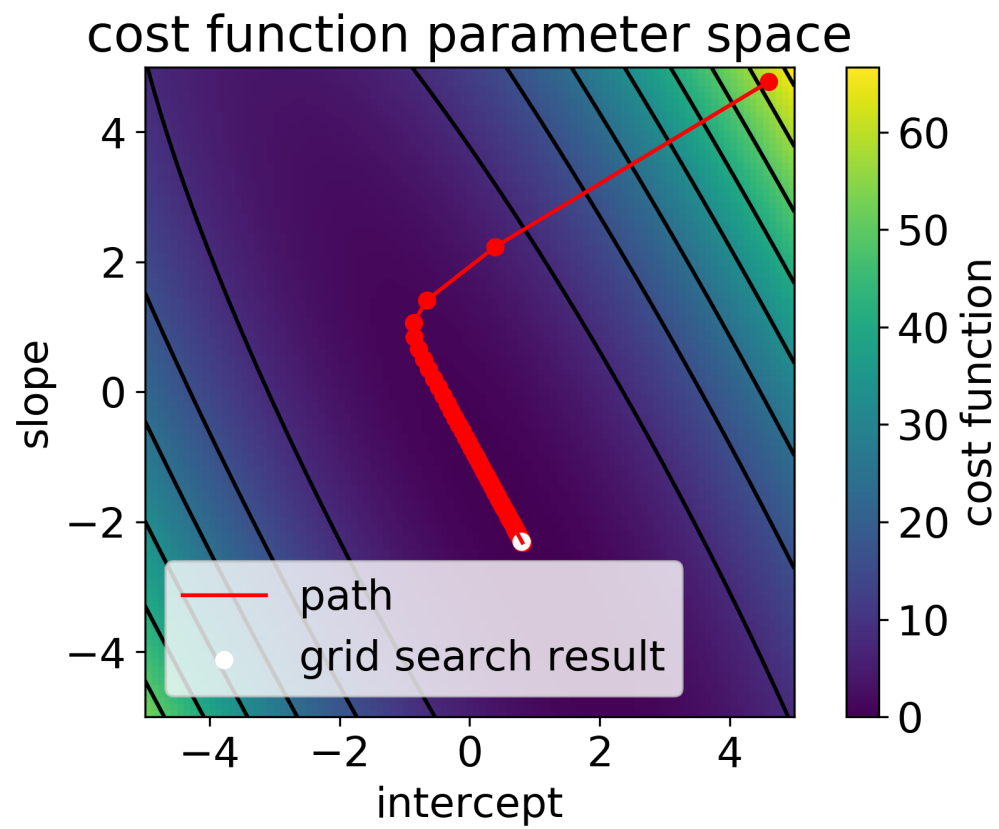
    return theta[0], cost_history, theta_history
```

```
In [15]: theta, cost_history, theta_hist = gradient_descent(df['x1'].values[:, np.newaxis], df['y'], [5.0, 5.0], 0.05, 1500)
print(theta)
print(theta_hist)
```

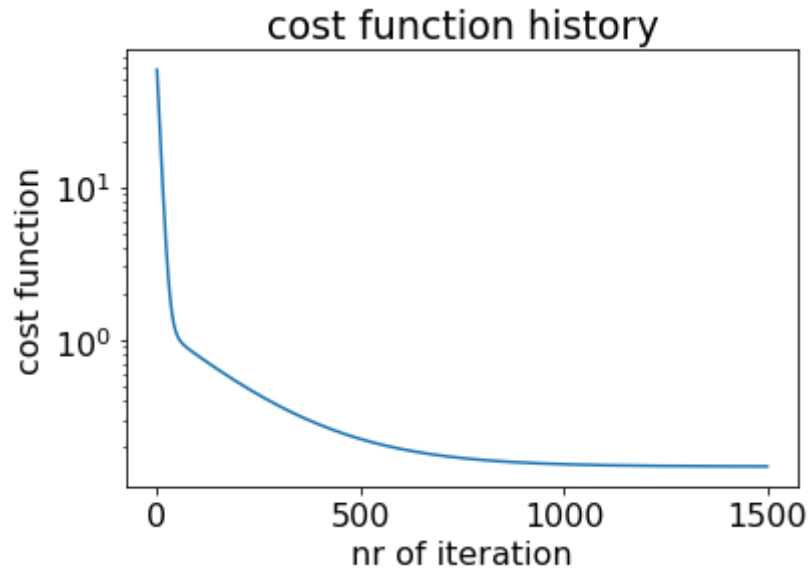
```
[ 0.81374506 -2.32315071]
[[ 4.60368338  4.77240776]
 [ 4.23299768  4.55868453]
 [ 3.88630699  4.35794734]
 ...
 [ 0.81353231 -2.3227565 ]
 [ 0.81363883 -2.32295387]
 [ 0.81374506 -2.32315071]]
```

```
In [16]: plt.figure(figsize=(6.4, 4.8))
ax = plt.gca()
extent = (np.min(theta0), np.max(theta0), np.min(theta1), np.max(theta1))
fig = ax.imshow(cost.T, origin='lower', extent=extent, vmin=0)
plt.colorbar(fig, label='cost function')
ax.contour(theta0, theta1, cost.T, levels=10, colors='black')
plt.plot(theta_hist[:, 0], theta_hist[:, 1], color='r', label='path')
plt.scatter(theta_hist[:, 0], theta_hist[:, 1], c='r')
plt.scatter(theta0[min_coords[0]], theta1[min_coords[1]], c='w', label='grid search result')
ax.xaxis.set_ticks_position("bottom")
plt.legend()
plt.xlabel('intercept')
plt.ylabel('slope')
plt.title('cost function parameter space')
plt.tight_layout()
plt.savefig('figures/cost_function_with_path.png', dpi=300)
plt.show()
```





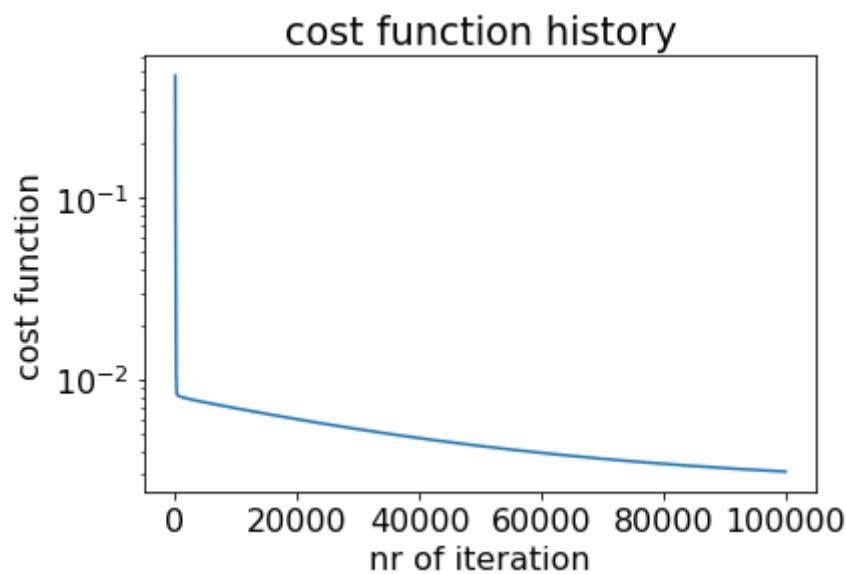
```
In [17]: plt.plot(cost_history)
plt.semilogy()
plt.ylabel('cost function')
plt.xlabel('nr of iteration')
plt.title('cost function history')
plt.savefig('figures/cost_hist.png',dpi=300)
plt.show()
```



```
In [18]: #theta,cost_history,theta_hist = gradient_descent(df[['x1','x2','x3','x4','x5']].values,df['y'],[1.1,-2.38,5.45,-44.7,75.33,-34.87],1.0,10000)
theta,cost_history,theta_hist = gradient_descent(df[['x1','x2','x3','x4','x5']].values,df['y'],[0,0,0,0,0,0],1.0,100000)
print(theta)

plt.plot(cost_history)
plt.semilogy()
plt.ylabel('cost function')
plt.xlabel('nr of iteration')
plt.title('cost function history')
plt.savefig('figures/cost_hist2.png',dpi=300)
plt.show()
```

```
[ 1.17809286 -1.46079625 -9.6913062   7.82929009  7.59160438 -5.4156485
 4]
```



By now you can

- Describe the most commonly used regression metrics
- Describe what the cost function is
- Explain how a simple gradient descent algorithm works

In []: